

## Web Security

When you think about adding users to your application, you typically have two aspects to consider:

1. **Authentication**—The process of creating users and letting them log in to your app  
The process of determining who you are.
2. **Authorization**—Customizing the experience and controlling what users can do, based on the current logged-in user.  
The process of determining what you're allowed to do.

Generally, you need to know who the user is before you can determine what they're allowed to do, so authentication always comes first, followed by authorization.

### 1 Understanding users and claims in ASP.NET Core

The concept of a user is baked in to ASP.NET Core. In chapter 3, you learned that the HTTP server, Kestrel, creates an `HttpContext` object for every request it receives. This object is responsible for storing all the details related to that request, such as the request URL, any headers sent, the body of the request, and so on.

The `HttpContext` object also exposes the current *principal* for a request as the `User` property. This is ASP.NET Core's view of which user made the request. Any time your app needs to know who the current user is, or what they're allowed to do, it can look at the `HttpContext.User` principal.

**DEFINITION** You can think of the *principal* as the user of your app.

In ASP.NET Core, principals are implemented as `ClaimsPrincipal`, which has a collection of *claims* associated with it, as shown in figure 14.1.

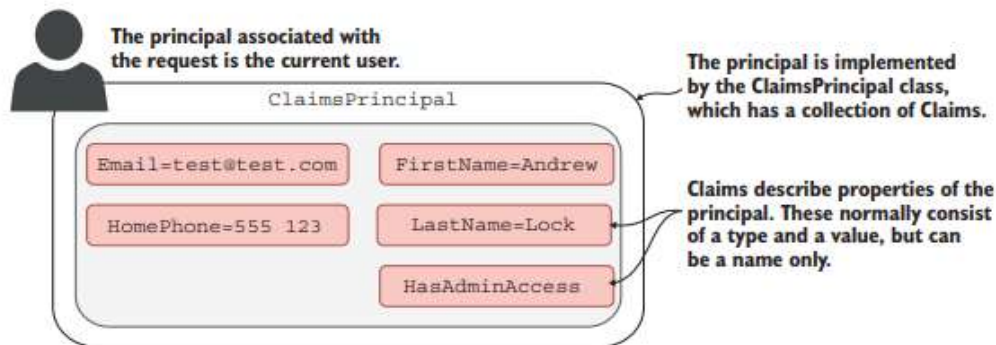


Figure 14.1 The principal is the current user, implemented as `ClaimsPrincipal`. It contains a collection of *Claims* that describe the user.

## SIGNING IN TO AN ASP.NET CORE APPLICATION

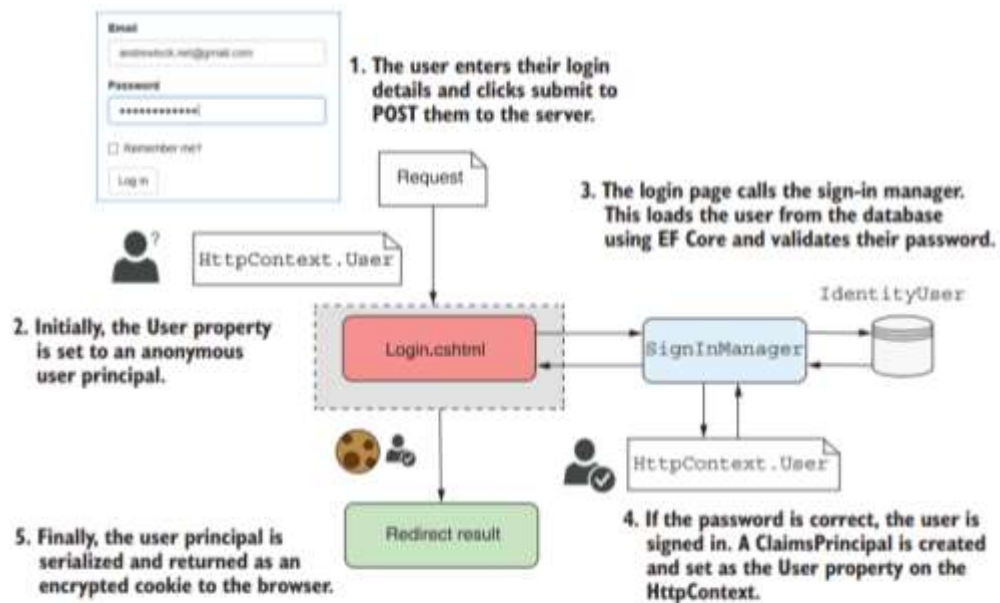


Figure 14.2 Signing in to an ASP.NET Core application. `SignInManager` is responsible for setting `HttpContext.User` to the new principal and serializing the principal to the encrypted cookie.

The meaty work happens inside the `SignInManager` service. This is responsible for loading a user entity with the provided username from the database and validating that the password they provided is correct.

**WARNING** Never store passwords in the database directly. They should be hashed using a strong one-way algorithm. The ASP.NET Core Identity system does this for you, but it's always wise to reiterate this point!

The key to persisting your identity across multiple requests lies in the final step of figure 14.2, where you serialized the principal in a cookie. Browsers will automatically send this cookie with all requests made to your app, so you don't need to provide your password with every request.

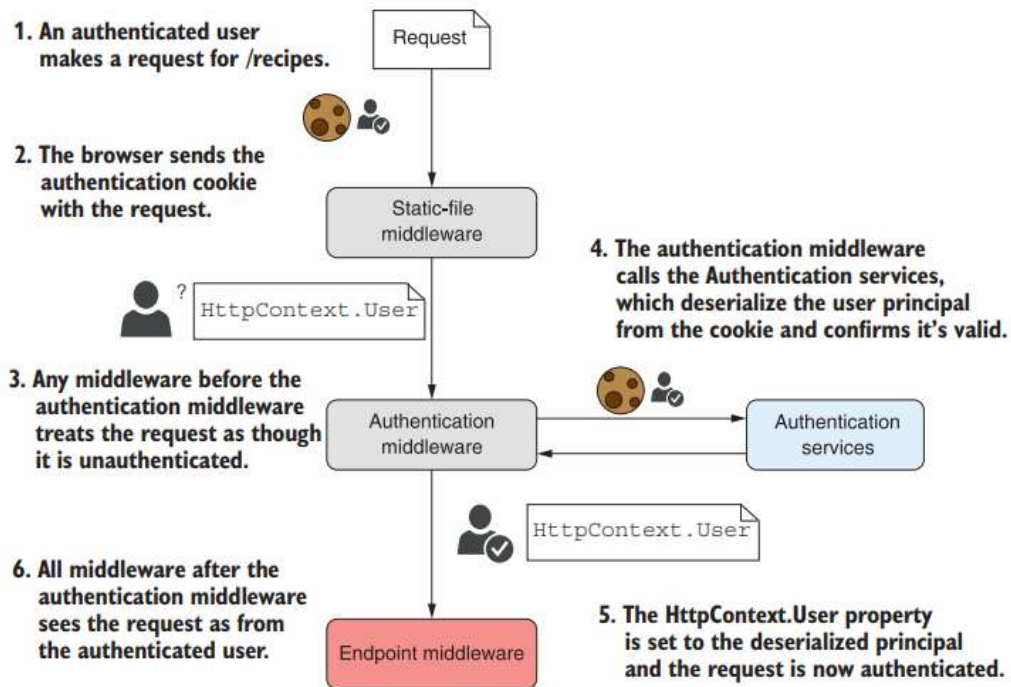
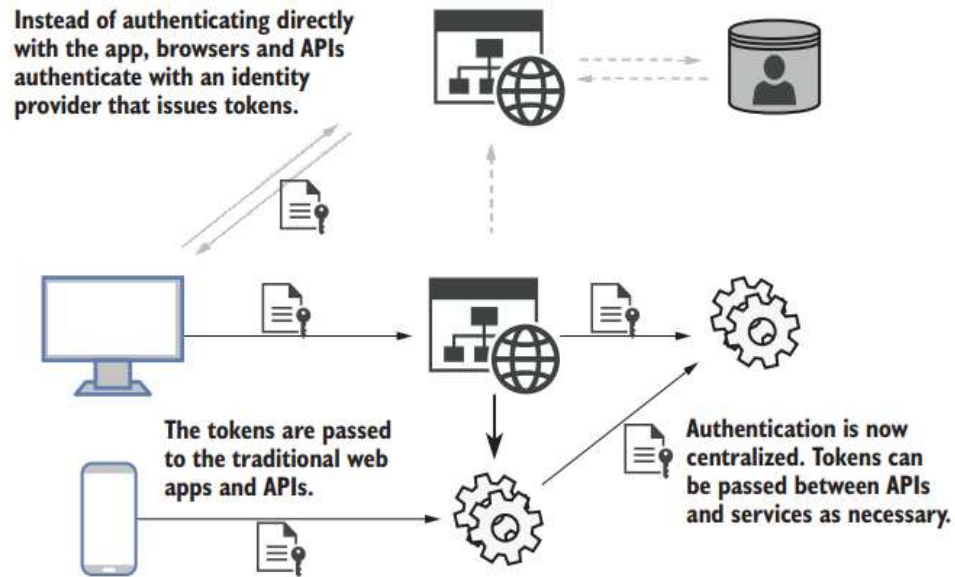


Figure 14.3 A subsequent request after signing in to an application. The cookie sent with the request contains the user principal, which is validated and used to authenticate the request.

**NOTE** The `AuthenticationMiddleware` is only responsible for authenticating incoming requests and setting the `ClaimsPrincipal` if the request contains an authentication cookie. It is not responsible for redirecting unauthenticated requests to the login page or rejecting unauthorized requests—that is handled by the `AuthorizationMiddleware`.



**Figure 14.6** An alternative architecture involves using a central identity provider to handle all the authentication and user management for the system. Tokens are passed back and forth between the identity provider, apps, and APIs.

This architecture is clearly more complicated on the face of it, as you’ve thrown a whole new service—the identity provider—into the mix, but in the long run this has a number of advantages:

- *Users can share their identity between multiple services.* As you’re logged in to the central identity provider, you’re essentially logged in to *all* apps that use that service. This gives you the single-sign-on experience, where you don’t have to keep logging in to multiple services.
- *Reduced duplication.* All of the sign-in logic is encapsulated in the identity provider, so you don’t need to add sign-in screens to all your apps.
- *Can easily add new providers.* Whether you use the identity provider approach or the traditional approach, it’s possible to use external services to handle the authentication of users. You’ll have seen this on apps that allow you to “log in using Facebook” or “log in using Google,” for example. If you use a centralized identity provider, adding support for additional providers can be handled in one place, instead of having to configure every app and API explicitly.

**TIP** Wherever possible, I recommend this approach, as it delegates security responsibilities to someone else. You can’t lose your user’s details if you never had them!

After creating ASP.net Core Web App with Identity services, following DI is created.

#### Listing 14.1 Adding ASP.NET Core Identity services to ConfigureServices

Adds the Identity system, including the default UI, and configures the user type as IdentityUser

ASP.NET Core Identity uses EF Core, so it includes the standard EF Core configuration.

```
public void ConfigureServices(IServiceCollection services)
{
    services.AddDbContext<ApplicationDbContext>(options =>
        options.UseSqlServer(
            Configuration.GetConnectionString("DefaultConnection")));
    services.AddDefaultIdentity<IdentityUser>(options =>
```

```
        options.SignIn.RequireConfirmedAccount = true)
        .AddEntityFrameworkStores<ApplicationDbContext>();
    services.AddRazorPages();
}
```

Configures Identity to store its data in EF Core

Requires users to confirm their accounts (typically by email) before they log in

The AddDefaultIdentity() extension method does several things:

- Adds the core ASP.NET Core Identity services.
- Configures the application user type to be IdentityUser. This is the entity model that is stored in the database and represents a "user" in your application. You can extend this type if you need to, but that's not always necessary, as you'll see in section 14.6.
- Adds the default UI Razor Pages for registering, logging in, and managing users.
- Configures token providers for generating 2FA and email confirmation tokens.



**Listing 14.2 Adding AuthenticationMiddleware to your middleware pipeline**

```
public void Configure(IApplicationBuilder app, IWebHostEnvironment env)
{
    app.UseStaticFiles();           ← Middleware placed before UseAuthentication
                                   ← will see all requests as anonymous.

    app.UseRouting();             ← The routing middleware determines which
                                   ← page is requested based on the request URL.

    app.UseAuthentication();       ← UseAuthentication should be
    app.UseAuthorization();         ← placed after UseRouting.

    app.UseEndpoints(endpoints =>
    {
        endpoints.MapRazorPages();
    });

    UseEndpoints should be last,
    after the user principal is set
    and authorization has been
    applied.

    UseAuthorization should be placed after
    UseAuthentication so it can access the user principal.
}
```

If you don't use this specific order of middleware, you can run into strange bugs where users aren't authenticated correctly, or where authorization policies aren't correctly applied. This order is configured for you automatically in your templates, but

**You need to update the database for Identity services using Nuget Package Console: update-database.**

**This will create all the necessary Database and Tables for storing user information, explore the database using Server Explorer menu.**

**It's important to understand the difference between the IdentityUser entity (stored in theAspNetUsers table) and the ClaimsPrincipal, which is exposed on HttpContext.User. When a user first logs in, an IdentityUser is loaded from the database. This entity is combined with additional claims for the user from the AspNetUserClaims table to create a ClaimsPrincipal. It's this ClaimsPrincipal that is used for authentication and is serialized to the authentication cookie, not the IdentityUser.**

**It's useful to have a mental model of the underlying database schema Identity uses, but in day-to-day work, you shouldn't have to interact with it directly—that's what Identity is for, after all! In the next section, we'll look at the other end of the scale— the UI of the app, and what you get out of the box with the default UI.**

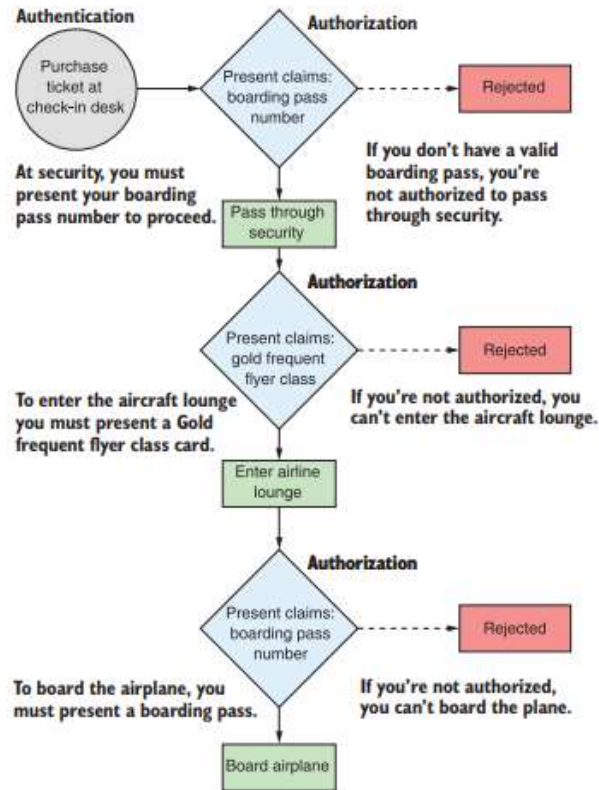


Figure 15.1 When boarding a plane at an airport, you pass through several authorization steps. At each authorization step, you must present a claim in the form of a boarding pass or a frequent flyer card. If you're not authorized, access will be denied.

See the Airport Project Code for the implementation of this process.

#### Listing 15.7 Creating an authorization policy with multiple requirements

```

public void ConfigureServices(IServiceCollection services)
{
    services.AddAuthorization(options =>
    {
        options.AddPolicy(
            "CanEnterSecurity",
            policyBuilder => policyBuilder
                .RequireClaim(Claims.BoardingPassNumber));
        options.AddPolicy(
            "CanAccessLounge",
            policyBuilder => policyBuilder.AddRequirements(
                new MinimumAgeRequirement(18),
                new AllowedInLoungeRequirement()));
    });
    // Additional service configuration
}
  
```

Adds a new policy for the airport lounge, called CanAccessLounge

Adds the previous simple policy for passing through security

Adds an instance of each IAuthorizationRequirement object

You now have a policy called "CanAccessLounge" with two requirements, so you can apply it to a Razor Page or action method using the `[Authorize]` attribute, in exactly the same way you did for the "CanEnterSecurity" policy:

```

[Authorize("CanAccessLounge")]
public class AirportLoungeModel : PageModel
{
  }
  
```

When a request is routed to the `AirportLounge.cshtml` Razor Page, the authorize middleware executes the authorization policy and each of the requirements is inspected. But you saw earlier that the requirements are purely data; they indicate what needs to be fulfilled, but they don't describe how that has to happen. For that, you need to write some handlers.

#### CREATING AUTHORIZATION HANDLERS TO SATISFY YOUR REQUIREMENTS

Authorization handlers contain the logic of how a specific `IAuthorizationRequirement` can be satisfied. When executed, a handler can do one of three things:

- Mark the requirement handling as a success
- Not do anything
- Explicitly fail the requirement

Handlers should implement `AuthorizationHandler<T>`, where `T` is the type of requirement they handle. For example, the following listing shows a handler for `AllowedInLoungeRequirement` that checks whether the user has a claim called `FrequentFlyerClass` with a value of `Gold`.

**Listing 15.8** `FrequentFlyerHandler` for `AllowedInLoungeRequirement`

```
You must override the abstract HandleRequirementAsync method.

public class FrequentFlyerHandler :
    AuthorizationHandler<AllowedInLoungeRequirement>
{
    protected override Task HandleRequirementAsync(
        AuthorizationHandlerContext context,
        AllowedInLoungeRequirement requirement)
    {
        if (context.User.HasClaim("FrequentFlyerClass", "Gold"))
        {
            context.Succeed(requirement);
        }
        return Task.CompletedTask;
    }
}
```

The handler implements `AuthorizationHandler<T>`.

The requirement instance to handle

The context contains details such as the `ClaimsPrincipal` user object.

Checks whether the user has the `FrequentFlyerClass` claim with the `Gold` value

If the user had the necessary claim, then mark the requirement as satisfied by calling `Succeed`.

If the requirement wasn't satisfied, do nothing.

This handler is functionally equivalent to the simple `RequireClaim()` handler you saw at the start of section 15.4, but using the requirement and handler approach instead.

When a request is routed to the `AirportLounge.cshtml` Razor Page, the authorization middleware sees the `[Authorize]` attribute on the endpoint with the `"CanAccess-Lounge"` policy. It loops through all the requirements in the policy, and all the handlers for each requirement, calling the `HandleRequirementAsync` method for each.

The authorization middleware passes the current `AuthorizationHandlerContext` and the requirement to be checked to each handler. The current `ClaimsPrincipal`



being authorized is exposed on the context as the `User` property. In listing 15.8, `FrequentFlyerHandler` uses the context to check for a claim called `FrequentFlyerClass` with the `Gold` value, and if it exists, indicates that the user is allowed to enter the airline lounge by calling `Succeed()`.

**NOTE** Handlers mark a requirement as being successfully satisfied by calling `context.Succeed()` and passing the requirement as an argument.

It's important to note the behavior when the user *doesn't* have the claim. `FrequentFlyerHandler` doesn't do anything if this is the case (it returns a completed `Task` to satisfy the method signature).

**NOTE** Remember, if any of the handlers associated with a requirement pass, then the requirement is a success. Only *one* of the handlers must succeed for the requirement to be satisfied.

This behavior, whereby you either call `context.Succeed()` or do nothing, is typical for authorization handlers. The following listing shows the implementation of `IsAirlineEmployeeHandler`, which uses a similar claim check to determine whether the requirement is satisfied.

#### Listing 15.9 `IsAirlineEmployeeHandler`

```
public class IsAirlineEmployeeHandler :
    AuthorizationHandler<AllowedInLoungeRequirement>
{
    protected override Task HandleRequirementAsync(
        AuthorizationHandlerContext context,
        AllowedInLoungeRequirement requirement)
    {
        if (context.User.HasClaim(c => c.Type == "EmployeeNumber"))
        {
            context.Succeed(requirement);
        }
        return Task.CompletedTask;
    }
}
```

The handler implements `AuthorizationHandler<T>`.

You must override the abstract `HandleRequirementAsync` method.

If the user has the necessary claim, mark the requirement as satisfied by calling `Succeed`.

Checks whether the user has the `EmployeeNumber` claim

If the requirement wasn't satisfied, do nothing.

**TIP** It's possible to write very generic handlers that can be used with multiple requirements, but I suggest sticking to handling a single requirement only. If you need to extract some common functionality, move it to an external service and call that from both handlers.

This pattern of authorization handler is common,<sup>1</sup> but in some cases, instead of checking for a *success* condition, you might want to check for a *failure* condition. In the

<sup>1</sup> I'll leave the implementation of `MinimumAgeHandler` for `MinimumAgeRequirement` as an exercise. You can find an example in the code samples for the chapter.

airport example, you don't want to authorize someone who was previously banned from the lounge, even if they would otherwise be allowed to enter.

You can handle this scenario by using the `context.Fail()` method exposed on the context, as shown in the following listing. Calling `Fail()` in a handler will always cause the requirement, and hence the whole policy, to fail. You should only use it when you want to guarantee failure, even if other handlers indicate success.

**Listing 15.10** Calling `context.Fail()` in a handler to fail the requirement

```
public class BannedFromLoungeHandler :
    AuthorizationHandler<AllowedInLoungeRequirement>
{
    protected override Task HandleRequirementAsync(
        AuthorizationHandlerContext context,
        AllowedInLoungeRequirement requirement)
    {
        if (context.User.HasClaim(c => c.Type == "IsBanned"))
        {
            context.Fail();
        }

        return Task.CompletedTask;
    }
}
```

The handler implements `AuthorizationHandler<T>`.

You must override the abstract `HandleRequirementAsync` method.

Checks whether the user has the `IsBanned` claim

If the user has the claim, fail the requirement by calling `Fail`. The whole policy will fail.

If the claim wasn't found, do nothing.

In most cases, your handlers will either call `Succeed()` or will do nothing, but the `Fail()` method is useful when you need a kill-switch to guarantee that a requirement won't be satisfied.

**NOTE** Whether a handler calls `Succeed()`, `Fail()`, or neither, the authorization system will always execute all of the handlers for a requirement, and all the requirements for a policy, so you can be sure your handlers will always be called.

The final step to complete your authorization implementation for the app is to register the authorization handlers with the DI container, as shown in the following listing.

**Listing 15.11** Registering the authorization handlers with the DI container

```
public void ConfigureServices(IServiceCollection services)
{
    services.AddAuthorization(options =>
    {
        options.AddPolicy(
            "CanEnterSecurity",
            policyBuilder => policyBuilder
                .RequireClaim(Claims.BoardingPassNumber));
        options.AddPolicy(
            "CanAccessLounge",
```

```

        policyBuilder => policyBuilder.AddRequirements(
            new MinimumAgeRequirement(18),
            new AllowedInLoungeRequirement()
        );
    };
    services.AddSingleton<IAuthorizationHandler, MinimumAgeHandler>();
    services.AddSingleton<IAuthorizationHandler, FrequentFlyerHandler>();
    services
        .AddSingleton<IAuthorizationHandler, BannedFromLoungeHandler>();
    services
        .AddSingleton<IAuthorizationHandler, IsAirlineEmployeeHandler>();
    // Additional service configuration
}

```

For this app, the handlers don't have any constructor dependencies, so I've registered them as singletons with the container. If your handlers have scoped or transient dependencies (the EF Core `DbContext`, for example), you might want to register them as scoped instead, as appropriate.

**NOTE** Services are registered with a lifetime of either transient, scoped, or singleton, as discussed in chapter 10.

You can combine the concepts of policies, requirements, and handlers in many ways to achieve your goals for authorization in your application. The example in this section, although contrived, demonstrates each of the components you need to apply authorization declaratively at the action method or Razor Page level, by creating policies and applying the `[Authorize]` attribute as appropriate.

As well as applying the `[Authorize]` attribute explicitly to actions and Razor Pages, you can also configure it globally, so that a policy is applied to every Razor Page or controller in your application. Additionally, for Razor Pages you can apply different authorization policies to different folders. You can read more about applying authorization policies using conventions in Microsoft's "Razor Pages authorization conventions in ASP.NET Core" documentation: <http://mng.bz/nMm2>.

There's one area, however, where the `[Authorize]` attribute falls short: resource-based authorization. The `[Authorize]` attribute attaches metadata to an endpoint, so the authorization middleware can authorize the user *before* an endpoint is executed, but what if you need to authorize the action *during* the action method or Razor Page handler?

This is common when you're applying authorization at the document or resource level. If users are only allowed to edit documents they created, then you need to load the document before you can tell whether they're allowed to edit it! This isn't easy with the declarative `[Authorize]` attribute approach, so you must use an alternative, imperative approach. In the next section, you'll see how to apply this resource-based authorization in a Razor Page handler.